

# Twitter Sentiment Analysis

User Manual

Version 1.0

April 30, 2025

MLOPS AI PROJECT

# Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	About This Application . . . . .	3
1.2	System Requirements . . . . .	3
1.2.1	Minimum Requirements . . . . .	3
1.2.2	Recommended . . . . .	3
<b>2</b>	<b>Getting Started</b>	<b>4</b>
2.1	Installation . . . . .	4
2.1.1	Step 1: Install Docker . . . . .	4
2.1.2	Step 2: Download the Application . . . . .	4
2.1.3	Step 3: Launch the Application . . . . .	4
2.1.4	Step 4: Verify Installation . . . . .	4
<b>3</b>	<b>Using the Application</b>	<b>6</b>
3.1	The Interface . . . . .	6
3.2	Analyzing Text . . . . .	7
3.2.1	Single Text Analysis . . . . .	7
3.2.2	Batch Analysis(Future Tasks) . . . . .	7
3.3	Understanding Results . . . . .	7
<b>4</b>	<b>Dashboard &amp; Analytics</b>	<b>8</b>
4.1	Accessing the Dashboard . . . . .	8
4.2	Dashboard Panels . . . . .	8
4.3	Twitter Sentiment Dashboard . . . . .	8
4.3.1	System Performance Metrics . . . . .	8
4.3.2	Application Performance . . . . .	9
4.3.3	Dashboard Configuration . . . . .	9
4.4	Grafana Interface Elements . . . . .	9
4.5	Customizing Your View . . . . .	9
4.6	Key Observations . . . . .	10
<b>5</b>	<b>Maintenance</b>	<b>11</b>
5.1	Regular Maintenance Tasks . . . . .	11
5.1.1	Daily Checks . . . . .	11
5.1.2	Weekly Tasks . . . . .	11
5.1.3	Monthly Activities . . . . .	11
5.2	Updating the System . . . . .	11

<b>6</b>	<b>Evaluation &amp; Testing</b>	<b>12</b>
6.1	Model Evaluation . . . . .	12
6.1.1	Performance Metrics . . . . .	12
6.1.2	Classification Report . . . . .	12
6.2	Model Monitoring . . . . .	13
6.2.1	Drift Detection . . . . .	13
6.2.2	Performance Over Time . . . . .	13
6.3	Unit Testing . . . . .	13
6.3.1	Test Types . . . . .	13
6.3.2	Running Tests . . . . .	13
6.4	Test Cases . . . . .	13
6.4.1	Data Pipeline Tests . . . . .	13
6.4.2	Model Tests . . . . .	14
6.4.3	Example Test Case . . . . .	14
<b>7</b>	<b>Troubleshooting</b>	<b>15</b>
7.1	Common Issues . . . . .	15
<b>8</b>	<b>Appendix</b>	<b>16</b>
8.1	Technical Details . . . . .	16
8.1.1	System Architecture . . . . .	16
8.1.2	File Structure . . . . .	16
8.2	Command Reference . . . . .	17
8.3	Glossary . . . . .	17

# Chapter 1

## Introduction

### 1.1 About This Application

The Twitter Sentiment Analysis application helps you understand the emotional tone behind tweets. Using advanced machine learning, it classifies text as positive or negative, allowing you to:

- Monitor brand perception
- Track campaign effectiveness
- Analyze customer feedback

### 1.2 System Requirements

#### 1.2.1 Minimum Requirements

- Operating System: Windows 10/11, macOS 10.15+, or Linux Ubuntu 20.04+
- Docker Desktop (latest version)
- 4GB RAM
- 10GB free disk space
- Internet connection

#### 1.2.2 Recommended

- 8GB+ RAM
- SSD storage
- Broadband internet connection

# Chapter 2

## Getting Started

### 2.1 Installation

Follow these simple steps to install the Twitter Sentiment Analysis application:

#### 2.1.1 Step 1: Install Docker

If you don't have Docker installed:

1. Visit [Docker's official website](#)
2. Download and install the version for your operating system
3. Launch Docker Desktop and ensure it's running

#### 2.1.2 Step 2: Download the Application

1. Download the application package from the provided link or clone from GitHub:

```
1 git clone https://github.com/Rajnishmaurya/twitter-sentiment-AI-project
```

2. Navigate to the application folder:

```
1 cd twitter-sentiment-AI-project
```

#### 2.1.3 Step 3: Launch the Application

1. Open a terminal or command prompt
2. Run the following command:

```
1 docker-compose up --build -d
```

3. Wait for the setup to complete (approximately 5-7 minutes)

#### 2.1.4 Step 4: Verify Installation

After installation, you can access:

- Main application: <http://localhost:8000>
- Monitoring dashboard: <http://localhost:3000>

- Default login: **admin**
- Default password: **admin** (please change after first login)

## Chapter 3

# Using the Application

### 3.1 The Interface

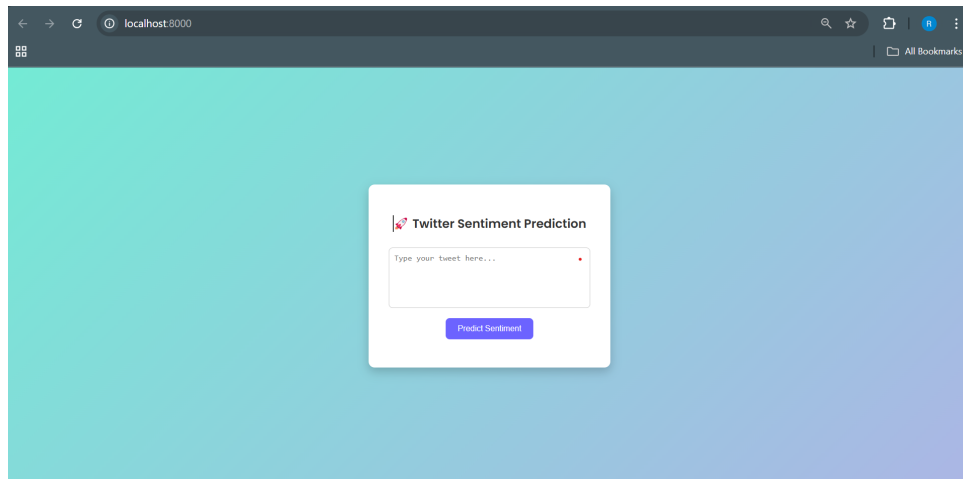


Figure 3.1: Twitter Sentiment Analysis Interface

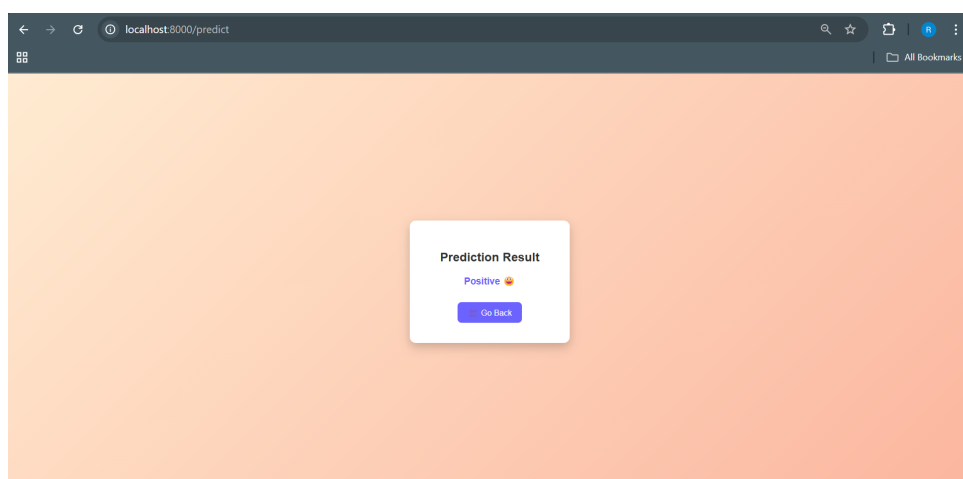


Figure 3.2: Twitter Sentiment Analysis Predict

1. **Input Area** - Enter tweets or text you want to analyze
2. **Predict Button** - Click to analyze the entered text

3. **Results Display** - Shows sentiment classification and confidence score
4. **History Panel** - Lists your recent analyses

## 3.2 Analyzing Text

### 3.2.1 Single Text Analysis

1. Enter text in the input field
2. Click "Predict Sentiment"
3. View the results:
  - **Positive** results appear in green with a 😊 emoji
  - **Negative** results appear in red with a ☹️ emoji

### 3.2.2 Batch Analysis(Future Tasks)

For analyzing multiple tweets:

1. Prepare a CSV file with a column named "text" containing your tweets
2. Click the "Batch Analysis" tab
3. Upload your CSV file
4. Click "Analyze All"
5. Download the results when processing completes

## 3.3 Understanding Results

The application provides several insights:

- **Sentiment Classification** - Positive or Negative
- **Visualization** - Color-coded representation of sentiment



## Chapter 4

# Dashboard & Analytics

### 4.1 Accessing the Dashboard

1. Navigate to <http://localhost:3000> in your browser
2. Log in with your credentials (default: admin/admin)
3. Select the "Twitter Sentiment Dashboard" from the home screen

### 4.2 Dashboard Panels

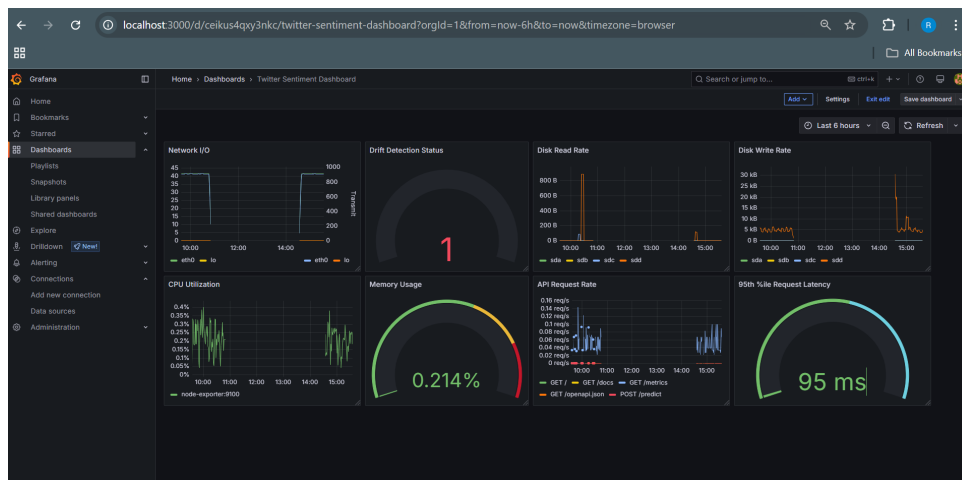


Figure 4.1: Dashboard Overview

The dashboard contains several informative panels:

### 4.3 Twitter Sentiment Dashboard

This dashboard displays monitoring metrics for a Twitter sentiment analysis system running in Grafana.

#### 4.3.1 System Performance Metrics

- **Network I/O:** Shows network throughput with values up to 45 on left axis and 1000 on right axis, displaying eth0 and lo interfaces

- **CPU Utilization:** Low CPU usage (0.05% to 0.4%) for node-exporter:9100
- **Memory Usage:** Currently at 0.214% utilization, displayed as a gauge visualization with green-yellow-red gradient
- **Disk Performance:**
  - Disk Read Rate: Shows periodic spikes up to approximately 800 B
  - Disk Write Rate: Shows periodic activity with values up to 30 kB

#### 4.3.2 Application Performance

- **API Request Rate:** Shows request rates for different endpoints (GET /, GET /docs, GET /metrics, GET /openapi.json, POST /predict)
- **95th %ile Request Latency:** Currently at 95 ms, displayed as a gauge visualization
- **Drift Detection Status:** Shows a value of 1 (likely indicating normal operation)

#### 4.3.3 Dashboard Configuration

- **Time Range:** Last 6 hours selected in the dropdown
- **Refresh Rate:** Configurable via dropdown
- **Dashboard Controls:** Options to Add panels, Settings, Exit edit mode, and Save dashboard

### 4.4 Grafana Interface Elements

- **Navigation:** Left sidebar with Home, Bookmarks, Starred, Dashboards sections
- **Dashboards Section:** Expandable subsections for Playlists, Snapshots, Library panels, Shared dashboards
- **Additional Features:** Drilldown, Alerting, Connections, Administration sections available

### 4.5 Customizing Your View

To customize your dashboard:

1. Click the "Add" button in the top right to add new panels
2. Use the "Settings" button to modify dashboard properties
3. Click directly on a panel and select "Edit" to modify specific visualizations
4. Adjust the time range using the "Last 6 hours" dropdown
5. Use the refresh button to update data manually or set an auto-refresh interval

## 4.6 Key Observations

- The system appears to be running with minimal resource utilization (0.214% memory, low CPU)
- API endpoints show variable request rates with clear patterns
- The 95ms latency indicates good performance for the service
- Storage I/O shows periodic activity patterns rather than sustained load
- This appears to be a monitoring dashboard for a Twitter sentiment analysis API rather than displaying the sentiment analysis results themselves

# Chapter 5

## Maintenance

### 5.1 Regular Maintenance Tasks

#### 5.1.1 Daily Checks

- Monitor dashboard for unusual activity
- Verify the application is accessible

#### 5.1.2 Weekly Tasks

- Review log files for warnings or errors
- Check disk space usage

#### 5.1.3 Monthly Activities

- Review model performance metrics
- Update Docker images if newer versions are available

### 5.2 Updating the System

When updates become available:

1. Download the latest version
2. Stop the current instance:
3. Replace application files with the new version
4. Rebuild and restart:

```
1 docker-compose down
```

```
1 docker-compose up --build -d
```

## Chapter 6

# Evaluation & Testing

### 6.1 Model Evaluation

#### 6.1.1 Performance Metrics

The sentiment analysis model is evaluated using the following metrics:

- **Accuracy:** Overall correctness of predictions (currently 75.9%)
- **Precision:** Proportion of positive identifications that are correct
- **Recall:** Proportion of actual positives that are identified correctly
- **F1 Score:** Harmonic mean of precision and recall (currently 0.76)

#### 6.1.2 Classification Report

Classification Report:				
	precision	recall	f1-score	support
	precision	recall	f1-score	support
0	0.78	0.73	0.75	159656
1	0.74	0.79	0.77	159547
accuracy			0.76	319203
	precision	recall	f1-score	support
0	0.78	0.73	0.75	159656
1	0.74	0.79	0.77	159547
	precision	recall	f1-score	support
0	0.78	0.73	0.75	159656
	precision	recall	f1-score	support
	precision	recall	f1-score	support
	precision	recall	f1-score	support
0	0.78	0.73	0.75	159656
accuracy			0.76	319203
macro avg	0.76	0.76	0.76	319203
weighted avg	0.76	0.76	0.76	319203

Figure 6.1: Classification Report for Sentiment Classification

## 6.2 Model Monitoring

### 6.2.1 Drift Detection

The system continuously monitors for concept drift:

- **Data Drift:** Changes in input data distribution
- **Concept Drift:** Changes in the relationship between input and output
- **Model Drift:** Degradation of model performance over time

### 6.2.2 Performance Over Time

The system tracks performance metrics over time to identify:

- Gradual degradation requiring retraining
- Sudden drops indicating potential issues
- Seasonal or cyclical performance patterns

## 6.3 Unit Testing

### 6.3.1 Test Types

- **Unit Tests:** Test individual functions and classes

### 6.3.2 Running Tests

To run the test suite:

```
1 # Run unit tests
2 pytest tests/
```

## 6.4 Test Cases

### 6.4.1 Data Pipeline Tests

#### Data Ingestion Tests

- Tests that data is loaded correctly from various sources
- Verifies handling of different file formats
- Checks error handling for corrupted data
- Validates schema enforcement

#### Data Validation Tests

- Tests for detection of invalid data
- Verifies data quality checks
- Validates handling of missing values
- Checks statistical profiling of datasets

### Preprocessing Tests

- Tests text normalization
- Verifies tokenization and stemming
- Validates feature extraction
- Checks data transformation pipelines

## 6.4.2 Model Tests

### Training Tests

- Tests model initialization
- Verifies training procedures
- Validates hyperparameter tuning
- Checks model serialization and loading

### Inference Tests

- Tests prediction accuracy
- Verifies handling of edge cases
- Validates performance under load
- Checks response formatting

## 6.4.3 Example Test Case

Below is an example of a test case from the test suite:

```
1 def test_model_prediction():
2     with open(MODEL_PATH, "rb") as f:
3         model = pickle.load(f)
4     with open(VECTORIZER_PATH, "rb") as f:
5         vectorizer = pickle.load(f)
6
7     text = "I am so happy!"
8     transformed = vectorizer.transform([text])
9     prediction = model.predict(transformed)
10
11     assert isinstance(prediction, np.ndarray) or isinstance(prediction,
12                                                                list)
13     assert len(prediction) == 1
14     assert prediction[0] in [0, 1] # assuming binary classification
```

# Chapter 7

## Troubleshooting

### 7.1 Common Issues

#### Application Not Accessible

**Issue:** Cannot access the application at <http://localhost:8000>

**Solution:**

1. Check if Docker is running
2. Verify all containers are up with `docker ps`
3. Check logs with `docker logs twitter-sentiment-fastapi`
4. Restart the application with `docker-compose restart`

#### Slow Performance

**Issue:** Analysis takes longer than expected

**Solution:**

1. Check system resources (CPU, memory)
2. Reduce batch size for large analyses
3. Clear application cache in Settings

#### Inaccurate Results

**Issue:** Sentiment predictions seem incorrect

**Solution:**

1. Ensure text is clear and in English
2. Check for sarcasm or ambiguity which may confuse the model
3. Report example to improve future versions



# Chapter 8

## Appendix

### 8.1 Technical Details

#### 8.1.1 System Architecture

The application uses a microservices architecture:

- **Frontend:** FastAPI web interface
- **Backend:** Python-based ML inference service
- **Monitoring:** Prometheus and Grafana
- **ML Pipeline:** DVC and MLflow

#### 8.1.2 File Structure

```
1 twitter-sentiment-AI-project/  
2     data/  
3         raw/  
4         processed/  
5     notebooks/  
6         eda.ipynb  
7     notebooks/  
8     src/  
9         data/  
10             data_ingestion.py  
11             data_validation.py  
12             preprocessing.py  
13         models/  
14             train.py  
15         evaluation/  
16             evaluate.py  
17         deployment/  
18             api.py  
19             docker/  
20                 Dockerfile  
21  
22     configs/  
23         config.yaml  
24     tests/  
25         test_data_pipeline.py  
26         test_model.py  
27     requirements.txt
```

```

28     .gitignore
29     dvc.yaml
30     docker-compose.yml
31     README.md

```

## 8.2 Command Reference

Command	Description
<code>docker-compose up -d</code>	Start all services
<code>docker-compose down</code>	Stop all services
<code>docker logs [container]</code>	View container logs
<code>dvc repro</code>	Rerun ML pipeline
<code>mlflow ui</code>	Launch experiment tracking UI

Table 8.1: Common Commands

## 8.3 Glossary

- **Sentiment Analysis:** Determining whether text expresses positive or negative
- **ML (Machine Learning):** Computer systems that learn from data
- **SVM:** Support Vector Machine, the model used for classification
- **TF-IDF:** Term Frequency-Inverse Document Frequency, a text feature extraction method
- **Data Drift:** When new data differs significantly from training data
- **Docker:** Container platform for running applications consistently
- **FastAPI:** Web framework for building APIs
- **Prometheus:** Monitoring and alerting toolkit
- **Grafana:** Analytics and monitoring platform
- **MLflow:** Platform for ML lifecycle management
- **DVC:** Data Version Control for ML projects